PromoProof: Master 3-Minute Walkthrough Teleprompter

Second-by-Second Video Recording Script & Screen Cue Guide

Dipanshu Raj • **Roll No: 230375**
IIT Kanpur • Kuku Technologies AI Generalist Submission
Live Demo: **promoproof-bvhz.onrender.com**

Recording Setup:	**Pacing & Delivery Rules:**
• **Browser:** Chrome at 100% zoom on https://promoproof-bvhz.onrender.com/#/ with clean tabs. • **Preload:** Click *Export state* → *Upload / Load* to import live-api-run-revalidated.json from reviewer outputs.	• **Word Count:** ~370 words (~2m 40s speaking time at 140 wpm, leaving 20s for UI transitions). • **Tone:** Confident, technical, authoritative. Treat P016 failure as an engineering triumph.

SECOND-BY-SECOND WALKTHROUGH & TELEPROMPTER

Timestamp	Screen & Mouse Cue	Word-for-Word Narration (Teleprompter Script)
00:00 – 00:25 (25 sec)	Start on **Promo Discovery**. Hover over stats: *50 episodes, 98k windows, 16 proposals, 1 pass*.	*"Hi everyone, I'm Dipanshu Raj, Roll Number 230375 from IIT Kanpur. This is **PromoProof** for Problem 1: Find the Promos.* *A mathematically valid timestamp does not make a good ad. Across 50 episodes of 'I Married a Broken Soldier', our engine scanned over 98,000 legal windows, yet only **one** cleared our strict semantic quality gate. PromoProof keeps that reality visible to editors instead of pretending every retrieved cut is ready to post."*
00:25 – 00:55 (30 sec)	Click **Load synthetic demo** → **Run analysis**. Click **Source & timeline** tab to show green bars.	*"Our core design principle is a strict partition of responsibility: **Code owns the physics; AI owns the narrative.*** *We never ask an LLM to guess timestamps—models hallucinate numbers. Instead, pure Python algorithms enforce dialogue safety: if an actor is speaking, a cut inside that interval is strictly illegal. The engine retrieves candidate windows between 20 and 60 seconds, dedupes overlapping cuts, and only feeds pre-filtered candidates to the reasoning model, cutting token input costs by over 86%."*
00:55 – 01:30 (35 sec)	Switch to real run. Select **Proposal P001**. Highlight Hook, Cliffhanger, and Why this cut.	*"Now looking at the real 50-episode season. Ingesting all 3,700 blocks took under one cent of compute on Groq.* *Our top-ranked recommendation is **P001 from Episode 20**, running 53.7 seconds. In the first three seconds, it opens on a visual power dynamic—Catherine raising a teacup with a cold, dismissive glare. It cuts precisely on an urgent, unresolved cliffhanger: 'If I stop, Seamus, the lung fever in the valley wins.' It works because a cold viewer instantly grasps the stakes without needing the prior 19 episodes."*
01:30 – 02:05 (35 sec)	Select **Proposal P016**. Point to red **FAILED GATE** and STITCH_GATE_FAILED.	*"The assignment brief asked: 'Where does your tool pick badly?' Here is an actual validator flaw we caught and resolved.* *In our initial live run, candidate **P016**—a composite cut—was erroneously passed by the LLM despite having a stitch legibility score of only 1 out of 4 and a blank open question. The model was simply too lenient.* *Instead of accepting it, I tightened the deterministic code gates in apply_reviews to enforce hard sub-score thresholds. Replaying the recorded outputs correctly failed P016 with STITCH_GATE_FAILED. Code policed the AI."*
02:05 – 02:30 (25 sec)	Click **Continuity desk**. Scroll through character entries (BLOND NOBLEMAN).	*"In microdramas, characters evolve across episodes, and generic labels like 'Blond Nobleman' recur in Episodes 2, 9, and 44 as completely different background actors. Automated alias merging would destroy storyline continuity. The Continuity Desk scopes identities per-episode as reviewer flags rather than unverified global merges. Furthermore, all show scripts are processed in volatile memory only—ensuring Zero Data Retention for unreleased IP."*
02:30 – 03:00 (30 sec)	Back to **The shortlist**. Type editor note, click **Accept cut**, click **Export state**.	*"Finally, the human editor stays firmly in control: they can reject cuts, add production notes, and export exact frame-level timestamps into CSV format for the video editing suite.* *Scaling this to 100 shows a week costs under **one dollar** in total AI inference. 46 unit tests pass, and our live Docker deployment is running on Render.* *Thank you for reviewing my submission, and I look forward to speaking with the Kuku team!"*